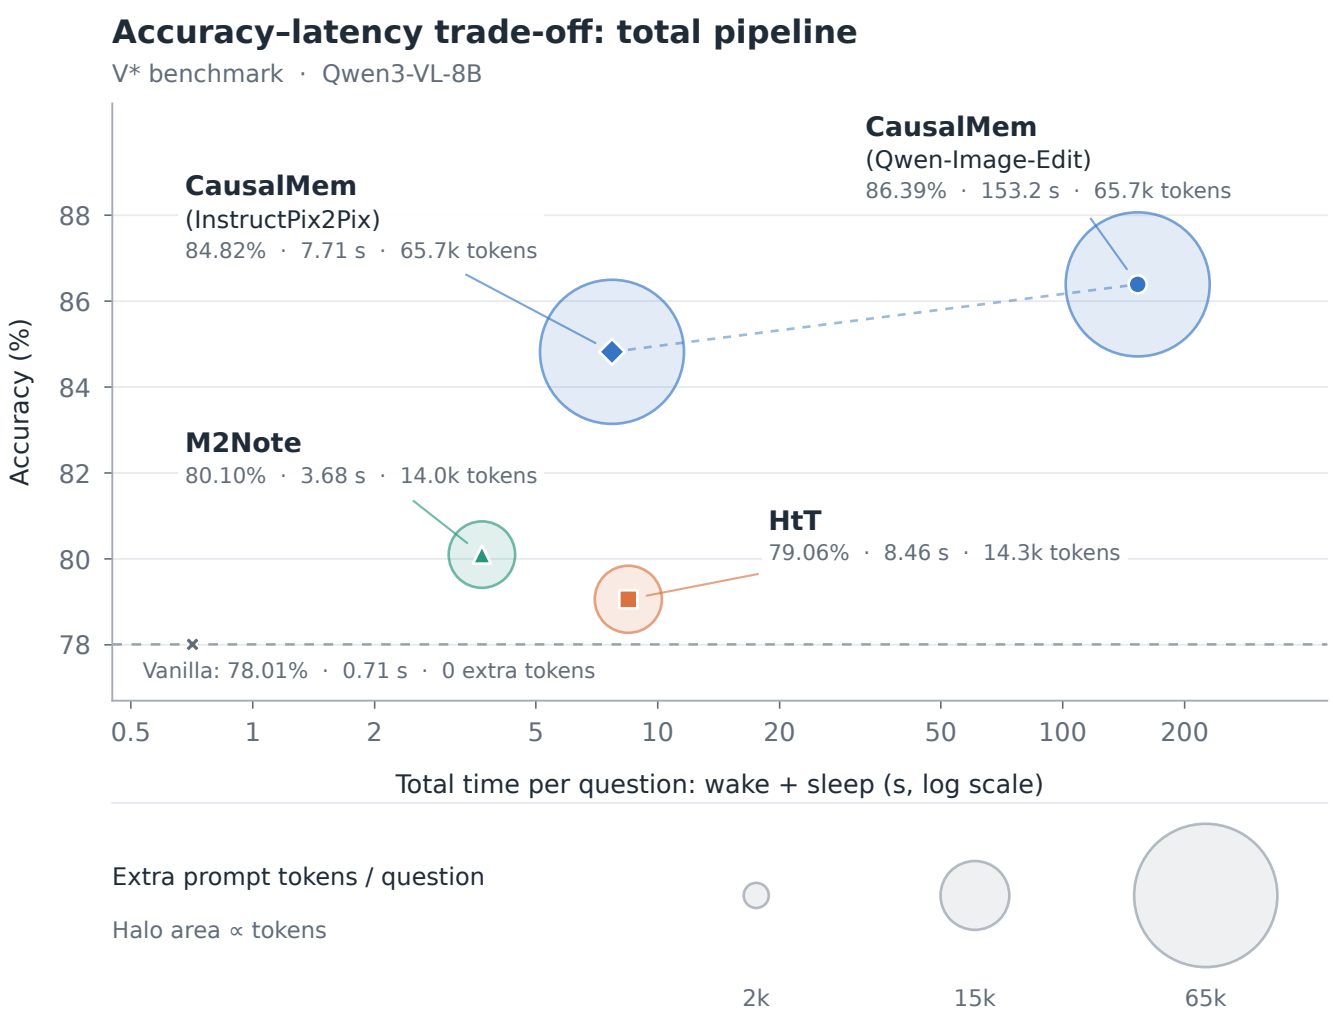
Accuracy–latency trade-off: total pipeline
V* benchmark · Qwen3-VL-8B
CausalMem
(Qwen-Image-Edit)
86.39% · 153.2 s · 65.7k tokens
CausalMem
(InstructPix2Pix)
84.82% · 7.71 s · 65.7k tokens
M2Note
80.10% · 3.68 s · 14.0k tokens
HtT
79.06% · 8.46 s · 14.3k tokens
Vanilla: 78.01% · 0.71 s · 0 extra tokens
Accuracy (%)
88
86
84
82
80
78
0.5
1
2
5
10
20
50
100
200
Total time per question: wake + sleep (s, log scale)
Extra prompt tokens / question
Halo area ∝ tokens
2k
15k
65k