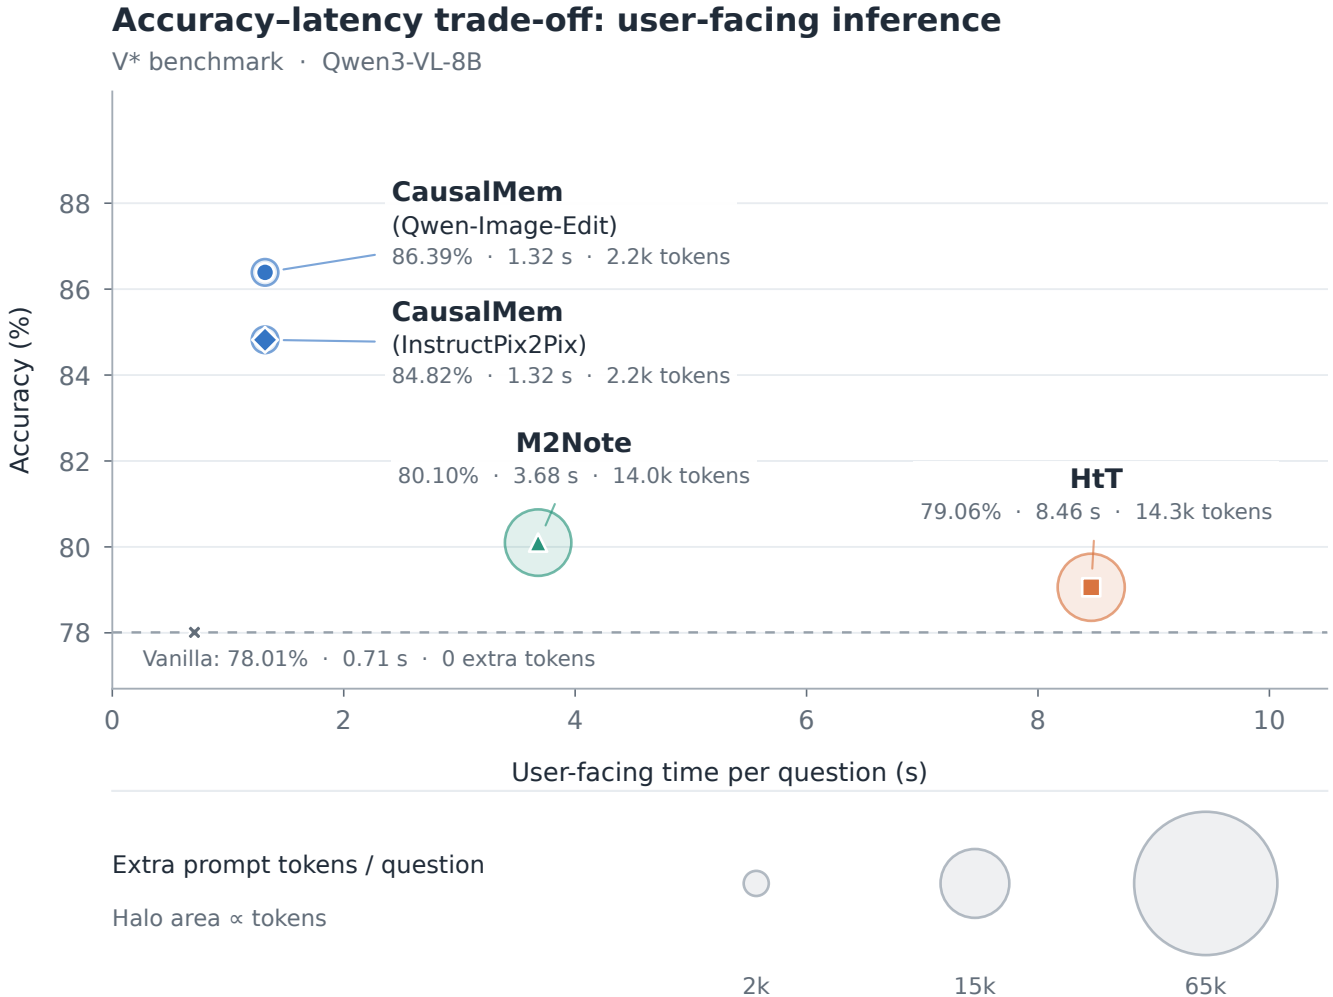
Accuracy–latency trade-off: user-facing inference
V* benchmark · Qwen3-VL-8B
CausalMem
(Qwen-Image-Edit)
86.39% · 1.32 s · 2.2k tokens
CausalMem
(InstructPix2Pix)
84.82% · 1.32 s · 2.2k tokens
M2Note
80.10% · 3.68 s · 14.0k tokens
HtT
79.06% · 8.46 s · 14.3k tokens
Vanilla: 78.01% · 0.71 s · 0 extra tokens
Accuracy (%)
88
86
84
82
80
78
0
2
4
6
8
10
User-facing time per question (s)
Extra prompt tokens / question
Halo area ∝ tokens
2k
15k
65k